

Music Recommender System

Team Name:

Flerken

Team members:

Xinghan Qin

Haoran Peng

Yueran Liu

Abstract

In this project, our team built a prediction model for the music recommender system. By giving more than 1 million user data including the score rating on artist, album, and genre of different tracks by Yahoo music users, we were asked to do data profiling first and then train our data to derive a model that's able to predict whether a specific like it or not.

To solve the task, we used Python in the entire process to load the data, make the dictionary to match all the features mentioned above to their track ID. After we got our sort data in a list in matrix form, we are finally able to apply four different classifiers by using the PySpark package. The results showed that we were able to make predictions about users' preferences with over 87% accuracy.

I.Data process

By giving more than 1 million user data including the score rating on artist, album, and genre of different tracks by Yahoo music users, we need to change the format which is convenient for us to handle.


Data set description

The data we used as following:

- albumData2.txt
- artistData2.txt
- genreData2.txt
- trackData2.txt
- trainItem2.txt
- testItem2.txt
- test2.txt

Feature addition

First of all, because our data is stored in a hierarchy structure, we decided to get some basic features based on this structure:



Then, to make a more accurate score and apply more data in the training data, we added more features:

1. Relative tracks: for each object track, we have one album. Hence, for the tracks under this album, we called them relative tracks. Then, we decided to use the maximum score, minimum score and mean score to represent all the relative tracks.
2. Relative albums: for each object track, we have one artist. Hence, for the album under this artist, we called them the relative album. Then, we decided to use the maximum score, minimum score and mean score to represent all the relative albums. However, for the tracks under all the relative albums, we decided not to count them because they were not very related to our object track.

To achieve this, we created the “albumData.txt” and “artistData.txt” by using “trackData2.txt”:

```
1 directory_name = "../../ee627a-2019fall"
2 testing_file_name = directory_name + "/testItem2.txt"
3 track_relationship_file_name = directory_name + "/trackData2.txt"
4
5 testing_file = open(testing_file_name, "r")
6 track_relationship_file = open(track_relationship_file_name, "r")
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000
```

```
import csv
output = open("albumData_final.txt", "w")
output.write("albumId, all trackId in this album\n")
for albumId in albumTrack.keys():
    trackList = albumTrack.get(albumId)
    output.write(albumId)
    for trackId in trackList:
        output.write(", " + trackId)
    output.write("\n")
```

```
albumId, all trackId in this album
106710,1,27943,95638,121434,126349,129878,129986,146801
280977,2,32098,37511,38304,72429,86624,128064,128978,142182,148087,164901,255111
38422,3,278462
119529,4,737,1057,22984,29439,48042,64537,99402,122002,148910,173813,193001,220350,222426
```

```
import csv
output = open("artistData_fianl.txt", "w")
output.write("artistId, all albumId in this album\n")
for artistId in artistTrack.keys():
    albumList = set(artistTrack.get(artistId))
    output.write(artistId)
    for albumId in albumList:
        output.write(", " + albumId)
    output.write("\n")
```

```
artistId, all albumId in this album
281667,237691,231576,106710,16320
233685,213241,180698,135761,152425,692,286126,106833,165758,130306,51795,66949,280977
219136,208450,38422,231573,142513
166863,92336,162148,22396,175518,153283,60110,150323,119529,121688,76204
294690,218479,167580,128171,147681,22766,200857,3460,112274,97383,180993,202045,16742,253490,835
44649,101746,68167,214870,120378,274309,59150
142773,181766,246442,141522,103130,83551,58498,184519
```

II.Track relation and score

To get our sorted data in a list in matrix form and fetch the related rating scores in the training data for the same users, we first read track files to make a dictionary. The key of the track dictionary is track id and the value of it is album, artist and genre id.

```
In [1]: 1 directory_name = "../ee627a-2019fall"
2 testing_file_name = directory_name + "/test2.txt"
3 track_file_name = directory_name + "/trackData2.txt"
4 testing_file = open(testing_file_name, "r")
5 tracking_file = open(track_file_name, "r")
```

```
In [2]: 1 print("Reading the track file ...")
2 dictrack = {}
3
4 for line in tracking_file:
5     arr_track = line.strip().split('|')
6     trackID = arr_track[0]
7     otherID = arr_track[1:]
8     trackdic={}
9     for others in otherID:
10         trackdic['album']=otherID[0]
11         trackdic['artist']=otherID[1]
12         trackdic['genre']=otherID[2:]
13     dictrack[trackID] = trackdic
14 print("Track file reading completed!")
```

```
Reading the track file ...
Track file reading completed!
```

```
In [3]: 1 dictrack
```

```
Out[3]: {'1': {'album': '106710',
               'artist': '281667',
               'genre': ['214765', '162234', '155788']},
         '2': {'album': '280977',
               'artist': '233685',
               'genre': ['131552', '173467', '48505']},
         '3': {'album': '38422',
               'artist': '219136',
               'genre': ['61215', '201738', '88853']},
         '4': {'album': '119529', 'artist': '166863', 'genre': ['17453', '35389']},
         '5': {'album': '16742',
               'artist': '294690',
               'genre': ['61215', '34486', '274088']},
         '7': {'album': '101746', 'artist': '44649', 'genre': ['198263']},
         '8': {'album': '58498',
               'artist': '142773',
               'genre': ['274161', '61215', '274088', '163949']},
         '9': {'album': 'None',
               'artist': 'None',
               'genre': ['163949', '163949', '163949']}}
```

With the dictionary above we wrote a list of track relation with features, each element of the list came up with: Userid, label, trackid, albumid, artistid, and all genreid

```
In [4]: 1 userdic={}
        2 scoredic={}
```

```
In [5]: 1 res = []
        2 for line in testing_file:
        3     arrTest = line.strip().split('|')
        4     userId=arrTest[0]
        5     trackId=arrTest[1]
        6     label=arrTest[2]
        7     scoredic[trackId]=arrTest[2]
        8     value=dictrack.get(trackId)
        9     res.append([userId,trackId, value.get('album'), value.get('artist'),value.get('genre'),label])
```

With this list, we wrote a new file of test relation. Each line of the file is the element in track relation list

```
In [7]: 1 output = open("test_relation.txt", "w")
        2 for i in res:
        3
        4     gen=i[4]
        5     gl=''
        6     for g in gen:
        7         gl=gl+' '+g
        8
        9     item=i[0]+' '+i[5]+' '+i[1]+' '+i[2]+' '+i[3]+gl
       10     output.write(item + "\n")
```

```
200031,1,30877,192723,132319,131552,176858,218185,251593,266073,67098
200031,1,8244,223220,233697,131552,176858,218185,75691,67098
200031,0,130183,None,None,139095,242383
200031,0,198762,220103,113265,131552,47898,201738,88853
200031,1,34503,43738,173170,131552,199606,51420,181006,67186,67331
200031,0,227283,183931,270460,131552,173467,110478,48505,133159
200032,0,218377,94189,178617,61215,34486,17453,35389
200032,0,110262,122655,52340,131552,176858,228419,218185
200032,1,18681,263622,211701,158282,242383,105160,110478,280261
200032,1,138493,58759,182642,173467,48505,133159
200032,1,64167,202904,46490,147886,161601
200032,0,22820,15669,103665,61215,78966,88853
200055,1,52198,209712,269659,77904
200055,0,233815,280805,205738,67098
200055,0,175557,194035,228477,199606,268304,172023,18161,88853
200055,0,59101,23848,159812,230141,201738
```

When we had all the features of these tracks, we used a training item to make the dictionary match all the features mentioned above to their track ID.

First of all, we read the test relation file and listed it.

[illegible]

Cause some track features did not have a score, we used 0 instead of None value. What's more, different tracks may have different kinds of genres. To fit all track, we choose three scores: genre max, genre min, and genre mean.

```

1 output = []
2 for users in test:
3     tmp = []
4     userId = users[0]
5     trackId = users[2]
6     want=users[1]
7     albumScore = res.get(userId).get(users[2])
8
9     if albumScore == None:
10         albumScore = "0"
11     artistScore = res.get(userId).get(users[3])
12     if artistScore == None:
13         artistScore = "0"
14     gen = []
15     if len(users[4]) == 0:
16         genMax = "0"
17         genMin = "0"
18         genMean = "0"
19     else:
20         for id in users[4]:
21             score = res.get(userId).get(id)
22             if score != None:
23                 gen.append(int(score))
24         if len(gen) != 0:
25             genMax = max(gen)
26             genMin = min(gen)
27             genMean = sum(gen) / len(gen)
28         else:
29             genMax = "0"
30             genMin = "0"
31             genMean = "0"
32     tmp = [userId,"",want,',', trackId,"", albumScore,"", artistScore,"", str(genMax),"
33     output.append(tmp)

```

The track id with all features(Userid, label, trackid, albumscore, artist score, genreamax, genreamin, genre mean)show as below, we used it in classifier:


```

test_021.txt
Userid,label,trackid,albumscore,artistscore,genreamax,genreamin,genreamean
200031,1,30877,0,90,90,50,73.33333333333333
200031,1,8244,0,90,90,80,85.0
200031,0,130183,0,0,0,0,0
200031,0,198762,0,0,90,90,90.0
200031,1,34503,0,90,90,50,73.33333333333333
200031,0,227283,0,0,90,50,76.66666666666667
200032,0,218377,0,0,0,0,0
200032,0,110262,0,0,0,0,0
200032,1,18681,0,90,90,90,90.0
200032,1,138493,0,90,90,90,90.0
200032,1,64167,0,90,90,90,90.0
200032,0,22820,0,0,0,0,0
200055,1,52198,0,90,90,90,90.0
200055,0,233815,0,0,0,0,0
200055,0,175557,0,0,0,0,0
200055,0,59101,0,0,0,0,0
200055,1,56695,0,90,90,90,90.0
200055,1,134398,0,0,90,90,90.0
200065,0,179571,0,0,0,0,0
200065,0,196286,0,0,0,0,0
200065,1,26875,0,90,90,90,90.0
200065,1,119451,0,0,90,90,90.0
200065,0,227397,0,0,0,0,0
200065,1,243829,0,0,70,70,70.0
200070,1,104207,0,0,0,0,0
200070,0,39252,0,0,90,80,85.0
200070,0,271459,0,0,0,0,0
200070,1,191512,0,70,90,90,90.0
200070,0,124239,0,0,0,0,0

```

We also used the same method to add more features, like the max score, min score and mean score of related tracks and albums.

```

test_score_feature_final01.txt
Userid,label,trackid,albumscore,trackMax,trackMin,trackMean,artistscore,albumMax,albumMin,albumMean,genreamax,genreamin,genreamean
199810,0,208019,0,0,0,0,0,0,0,0,0,0,0
199810,0,74139,0,0,0,0,0,0,0,0,80,80,80.0
199810,0,9903,0,0,0,0,0,0,0,0,0,0,0
199810,0,242681,0,0,0,0,0,0,0,0,0,0,0
199810,0,18515,0,0,0,0,70,0,0,0,0,0,0
199810,0,105760,0,0,0,0,90,0,0,0,80,80,80.0
199812,0,276940,0,0,0,0,0,0,0,0,0,0,0
199812,0,142408,100,0,0,0,100,100,100,100.0,80,80,80.0
199812,0,130023,100,0,0,0,100,100,100,100.0,80,80,80.0
199812,0,29189,0,0,0,0,0,0,0,0,80,80,80.0
199812,0,223706,0,0,0,0,100,100,100,100.0,80,80,80.0
199812,0,211361,0,0,0,0,0,0,0,0,0,0,0
199813,0,188441,0,0,0,0,90,0,0,0,80,80,80.0
199813,0,20968,0,0,0,0,0,0,0,0,80,80,80.0
199813,0,21571,90,0,0,0,90,90,90,90.0,0,0,0
199813,0,79640,0,0,0,0,90,90,90,90.0,80,80,80.0
199813,0,184173,0,0,0,0,70,0,0,0,80,80,80.0
199813,0,111874,0,0,0,0,0,0,0,0,80,80,80.0
199814,0,122375,0,0,0,0,0,0,0,0,0,0,0
199814,0,189043,75,76,76,76.0,75,75,75,75.0,0,0,0
199814,0,122429,0,0,0,0,0,0,0,0,0,0,0
199814,0,52519,0,0,0,0,0,0,0,0,0,0,0
199814,0,232332,100,0,0,0,100,100,100,100.0,0,0,0
199814,0,262193,75,76,76,76.0,75,75,75,75.0,0,0,0
199815,0,64345,0,0,0,0,50,70,70,70.0,80,50,65.0
199815,0,118841,0,0,0,0,0,0,0,0,0,0,0
199815,0,275682,100,100,30,75.0,100,100,50,91.25,80,80,80.0
199815,0,30062,0,90,30,55.714285714285715,100,100,100,100.0,80,60,73.33333333333333

```

III. Classifier

At the beginning, we didn't apply any model. We used our own experience to add weight to the features.

The first time, we ignored all the 'None' scores and set the weight of all features like 1. And then, according to the rank of the final score, we recommended the top 3 tracks.

```
my_track = {}
result = [{"TrackID", "Predictor"}]
count = 0

training_dict = read_training_file(training_file_name)
track_dict = read_track_file(track_relationship_file_name)

print("Start to predict rating ...")
for line in testing_file:
    arr_test = line.strip().split('|')
    if len(arr_test) == 2:
        user_id = arr_test[0]
        count = int(arr_test[1])
    else:
        track_id = arr_test[0]
        rate = predict_rate(user_id, track_id, training_dict, track_dict)
        my_track[track_id] = rate
        count = count - 1
    if count == 0:
        my_track_recommendation = predict_recommendation(my_track)
        for i in range(6):
            if i < 3:
                result.append([user_id + "_" + my_track_recommendation[i], 1])
            else:
                result.append([user_id + "_" + my_track_recommendation[i], 0])
        my_track = {}
print("Predict rating completed!")

print("Start to write down the prediction ...")
with open(output_file_name, "w") as output_file:
    writer = csv.writer(output_file)
    for item in result:
        writer.writerow(item)
print("Processing completed!")

def predict_rate(user_id, track_id, training_dict, track_dict):
    check_list = track_dict.get(track_id)
    score_dict = training_dict.get(user_id)
    count = 0
    res = 0
    for id in check_list:
        score = score_dict.get(id)
        if score != None:
            count = count + 1
            res = res + int(score)
    if count == 0:
        return 0
    res = res / count
    return res

def predict_recommendation(my_track):
    rates = list(my_track.values())
    tracks = list(my_track.keys())
    rates.sort(reverse=True)
    for item in tracks:
        r = my_track.get(item)
        i = rates.index(r)
        rates[i] = item
    return rates
```

The score was 0.79683.

The second time, we set the weight as 0.4 for the album, 0.3 for artist and 0.1 for each general:

```
def predict_rate(user_id, track_id, training_dict, track_dict):
    check_list = track_dict.get(track_id)
    score_dict = training_dict.get(user_id)
    count = 0
    res = 0
    kind = 0 # 0 for album, 1 for artist, 2 and more than 2 for genre
    just_genre = False
    for id in check_list:
        score = score_dict.get(id)
        if kind < 2:
            if score != None:
                if kind == 0:
                    count = count + 0.4
                    res = res + int(score) * 0.4
                else:
                    count = count + 0.3
                    res = res + int(score) * 0.3
            else:
                if kind == 2 and count == 0:
                    just_genre = True
                if score != None:
                    count = count + 0.1
                    res = res + int(score) * 0.1
        kind = kind + 1
    if count == 0:
        return 0
    res = res / count
    if just_genre:
        return res / 2
    return res
```

This time, the score had improvement and it was 0.85442.

Then, we applied four classifier model(Logistics regression model; DecisionTree; Random forest; GBTC) to the training data which had features below:

Userid,label,trackid,albumscore,artistscore,genreamax,genreamin,genreamean

First, we used the same method(VectorAssembler and PipeLine) to change the format of training data(test2, which has 6000 labeled data) to fit the model.

```
In [1]: from pyspark.sql import SparkSession
spark = SparkSession.builder.appName('GDBC').getOrCreate()

In [2]: df = spark.read.format("csv").option("sep", ",").option("inferSchema", "true").option("header", "true").load("test2_score.
df.printSchema()

root
|-- Userid: integer (nullable = true)
|-- label: integer (nullable = true)
|-- trackid: integer (nullable = true)
|-- albumscore: integer (nullable = true)
|-- artistscore: integer (nullable = true)
|-- genreamax: integer (nullable = true)
|-- genreamin: integer (nullable = true)
|-- genreamean: double (nullable = true)
```

```
In [3]: import pandas as pd
pd.DataFrame(df.take(5), columns=df.columns).transpose()
numeric_features = ['Userid', 'trackid', 'label', 'albumscore', 'artistscore', 'genreamax', 'genreamin', 'genreamean']
df.select(numeric_features).describe().toPandas().transpose()
```

```
Out[3]:
```

	0	1	2	3	4
summary	count	mean	stddev	min	max
Userid	6000	206286.491	3515.749299951101	200031	212234
trackid	6000	149260.0025	86146.45891885243	65	296098
label	6000	0.5	0.5000416718757232	0	1
albumscore	6000	0.0	0.0	0	0
artistscore	6000	23.488666666666667	38.63689523198974	0	100
genreamax	6000	48.80733333333333	43.19745227289881	0	100
genreamin	6000	42.272333333333336	40.789456944842996	0	100
genreamean	6000	45.58248076923077	41.17219269010034	0.0	100.0

```
In [4]: from pyspark.ml.feature import VectorAssembler
import pyspark.sql.types as types
df=df.withColumn('label',df['label'].cast(types.IntegerType()))
df=df.withColumn('Userid',df['Userid'].cast(types.IntegerType()))
df=df.withColumn('trackid',df['trackid'].cast(types.IntegerType()))
df=df.withColumn('albumscore',df['albumscore'].cast(types.IntegerType()))
df=df.withColumn('artistscore',df['artistscore'].cast(types.IntegerType()))
df=df.withColumn('genreamax',df['genreamax'].cast(types.IntegerType()))
df=df.withColumn('genreamin',df['genreamin'].cast(types.IntegerType()))
df=df.withColumn('genreamean',df['genreamean'].cast(types.IntegerType()))

assembler=VectorAssembler(inputCols=['Userid', 'trackid', 'label', 'artistscore', 'genreamax', 'genreamin', 'genreamean'],out
```

```
In [5]: from pyspark.ml import Pipeline
pipeline=Pipeline(stages=[assembler])
model=pipeline.fit(df)
train=model.transform(df)
```

```
In [6]: df1 = spark.read.format("csv").option("sep", ",").option("inferSchema", "true").option("header", "true").load("test_02.txt")
df1.printSchema()
```

```
root
|-- Userid: integer (nullable = true)
|-- label: integer (nullable = true)
|-- trackid: integer (nullable = true)
|-- albumscore: integer (nullable = true)
|-- artistscore: integer (nullable = true)
|-- genreamax: integer (nullable = true)
|-- genreamin: integer (nullable = true)
|-- genreamean: double (nullable = true)
```

```
In [7]: pd.DataFrame(df1.take(5), columns=df1.columns).transpose()
numeric_features = ['Userid', 'trackid', 'albumscore', 'artistscore', 'genreamax', 'genreamin', 'genreamean']
df1.select(numeric_features).describe().toPandas().transpose()
```

```
Out[7]:
```

	0	1	2	3	4
summary	count	mean	stddev	min	max
Userid	120000	224372.7714	14155.612277030812	199810	249010
trackid	120000	147871.92235833334	85421.1993529344	1	296099
albumscore	120000	23.685075	38.767846279190096	0	100
artistscore	120000	35.016175	42.73300263592536	0	100
genreamax	120000	35.471716666666666	40.80419698479321	0	100
genreamin	120000	32.74715833333333	39.44356509895082	0	100
genreamean	120000	34.149823609307354	39.70187105524825	0.0	100.0

```
In [8]: df1=df1.withColumn('userid',df1['userid'].cast(types.IntegerType()))
df1=df1.withColumn('label',df1['label'].cast(types.IntegerType()))
df1=df1.withColumn('trackid',df1['trackid'].cast(types.IntegerType()))
df1=df1.withColumn('albumscore',df1['albumscore'].cast(types.IntegerType()))
df1=df1.withColumn('artistscore',df1['artistscore'].cast(types.IntegerType()))
df1=df1.withColumn('genreamax',df1['genreamax'].cast(types.IntegerType()))
df1=df1.withColumn('genreamin',df1['genreamin'].cast(types.IntegerType()))
df1=df1.withColumn('genreamean',df1['genreamean'].cast(types.IntegerType()))

assembler1=VectorAssembler(inputCols=['userid','trackid','label','artistscore','genreamax','genreamin','genreamean'],ou

In [9]: pipeline=Pipeline(stages=[assembler1])
modell=pipeline.fit(df1)
test=modell.transform(df1)

In [10]: print("Training Dataset Count: " + str(train.count()))
print("Test Dataset Count: " + str(test.count()))

Training Dataset Count: 6000
Test Dataset Count: 120000
```

Second, we tried these four classifiers to generate different models. Then apply the models to test data(120000 unlabeled data).

```
from pyspark.ml.classification import DecisionTreeClassifier
dt = DecisionTreeClassifier(featuresCol = 'features', labelCol = 'label',maxDepth = 3,)
dtModel = dt.fit(train)
```

```
from pyspark.ml.classification import GBTClassifier
gbt = GBTClassifier(featuresCol = 'features',maxIter=10)
gbtModel = gbt.fit(train)
```

```
from pyspark.ml.classification import LogisticRegression
lr = LogisticRegression(featuresCol = 'features', maxIter=10,regParam=0.3, elasticNetParam=0.8)
lrModel = lr.fit(test)
```

```
from pyspark.ml.classification import DecisionTreeClassifier
dt = DecisionTreeClassifier(featuresCol = 'features', labelCol = 'label',maxDepth = 3,)
dtModel = dt.fit(test)
```

Finally, we extracted the prediction from the result. Note that something wrong appeared on the prediction (like more than three '0' or '1' prediction in one user) we transferred ordered probability to the final prediction

```
In [18]: pre=predictions1.rdd.map(lambda x: x.prediction).collect()
user=predictions1.rdd.map(lambda x: x.userid).collect()
track=predictions1.rdd.map(lambda x: x.trackid).collect()
```

```
In [19]: output = []
for i in range(120000):
    output.append([user[i],track[i],pre[i]])
```

```
In [20]: import csv
with open("GDBCOutput.csv", "w") as output_file:

    writer = csv.writer(output_file)
    writer.writerow(["TrackID", "Predictor"])
    for line in output:
        item = [str(line[0]) + " " + str(line[1]), line[2]]
        writer.writerow(item)
```

```
In [3]: file = open("res1.txt", "r")
res = {}
preId = "199810"
trackDic = {}
count = 0
userId = None
for line in file:
    arr_track = line.strip().split(',')
    userId = arr_track[0]
    trackId = arr_track[1]
    rate = arr_track[2]
    if userId != preId:
        res[preId] = trackDic
        trackDic = {}
        preId = userId
    trackDic[trackId] = float(rate)
    count += 1
res[userId] = trackDic
```

```
In [5]: result = []
for userId in res:
    my_track_recommendation = predict_recommendation(res.get(userId))
    for i in range(6):
        if i < 3:
            result.append([userId + "_" + my_track_recommendation[i][0], 0])
        else:
            result.append([userId + "_" + my_track_recommendation[i][0], 1])
```

```
In [7]: import csv
with open("res_ordered1.csv", "w") as output_file:
    writer = csv.writer(output_file)
    writer.writerow(["TrackID", "Predictor"])
    for line in result:
        writer.writerow(line)
```

Here are the scores of these Models

1. Logistics regression model(0.83256)
2. Decision tree(0.83859)
3. Random forest(0.74292)
4. GBTC(0.83552)

IV.Ensembling

We used the function below to obtain the weight of each submission.

$$\mathbf{a}_{LS} = (\mathbf{S}^T \mathbf{S})^{-1} \mathbf{S}^T \mathbf{x}$$

First, we changed all '0' in these data to '-1' to indicate the prediction and true are matched or not.

```
import pandas as pd
data = pd.read_csv('Submission_Ensembling0.81954.csv')
for i, row in data.iterrows():
    if data['Predictor'][i] == 0.0:
        data['Predictor'][i] = -1
data.to_csv('Submission_Ensembling0.81954_new.csv', index=False)
```

Second, we used six submissions to generate the ensemble matrix and score matrix.

```
# load all scores data from submission.csv
pre1=pd.read_csv('Submission_Ensembling0.81954.csv')
pre2=pd.read_csv('submission0.87005.csv')
pre3=pd.read_csv('LR_odered_0.83256.csv')
pre4=pd.read_csv('output0.85442.csv')
pre5=pd.read_csv('decisionTreeOutput_0.83859.csv')
pre6=pd.read_csv('GBTC_odered_0.80075.csv')
pre7=pd.read_csv('output_ALS0.66016.csv')
pre8=pd.read_csv('submission0.80925.csv')
pre9=pd.read_csv('submission0.84464.csv')
prenew=pd.read_csv('sample_submission.csv')

#list

s1=pre1.Predictor.tolist()
s2=pre2.Predictor.tolist()
s3=pre3.Predictor.tolist()
s4=pre4.Predictor.tolist()
s5=pre2.Predictor.tolist()
s6=pre3.Predictor.tolist()
s7=pre4.Predictor.tolist()
s8=pre3.Predictor.tolist()
s9=pre4.Predictor.tolist()
#S
S=np.transpose(mat([s1,s2,s3,s4,s5,s6,s7,s8,s9]))
S

matrix([[0, 0, 1, ..., 1, 1, 1],
        [1, 1, 1, ..., 1, 1, 1],
        [0, 0, 1, ..., 1, 1, 1],
        ...,
        [1, 1, 0, ..., 0, 0, 0],
        [0, 0, 0, ..., 0, 0, 0],
        [1, 1, 0, ..., 0, 0, 0]])

#score
#STy=k*(2*score_mat-1)
score1=0.81954
score2=0.87005
score3=0.83256
score4=0.85442
score5=0.83859
score6=0.80075
score7=0.66016
score8=0.80925
score9=0.84464
score_mat=np.transpose(mat([score1,score2,score3,score4,score5,score6,score7,score8,score9]))
k=120000
STy=k*(2*score_mat-1)
STy

matrix([[76689.6],
        [88812. ],
        [79814.4],
        [85060.8],
        [81261.6],
        [72180. ],
        [38438.4],
        [74220. ],
        [82713.6]])
```

Finally, we generated the function and applied it to the submission matrix then we had the ensembled submission.

```
#a=((STS)^(-1))*STy
STS=np.transpose(S)*S
STS_inverse = np.linalg.pinv(STS)
STS_inverse
a=STS_inverse*STy
#S0=S*a
S0=S*a
S0=S0.T
newS0=np.array(S0).reshape((20000,6)).tolist()
newS0

output = []
for i in range(120000):
    s = id1[i]
    arr_test = s.strip().split("_")
    string = arr_test[0] + "," + arr_test[1] + "," + str(S0.item(i))
    output.append(string)
```



```
file = open("res1.txt", "w")
for line in output:
    file.write(line + "\n")
```

With the file, we finally got the correct rate up to 0.87024.

6	Flerken		0.87024	55	13d
---	---------	--	---------	----	-----

Conclusion

In this project, we first used the basic method to predict the score in the test data by giving weight to each feature, but it lost some relation between each other. So we reformed the data by generating the relation between different user, track, album, artist and genre to a dictionary form. Then we applied four classifiers to these reformed data and ensemble the above submission together and got the final result(0.87024) and rank(top 8).

But we still have many problems that need to be solved. The classifiers we used are very common, and when we tried to obtain the weight of each submission, the accuracy of our submission did not improve in the beginning. So we have to keep learning in this area to find a better way to construct a more complicated and more accurate model to predict the result.

To get more information, please visit our Github:

https://github.com/OseankingO/627_project.git